Fusion by example

Joel de Guzman and Dan Marsden

Outline

- What is Fusion?
- A quick outline of Fusion
- 4 examples

What is Fusion?

- A library of heterogeneous containers and associated algorithms
- A fusion of compile time and run time programming techniques
- A tuple library

Tuples in the C++ world

- Conventional C++
 - std::pair
- C++ 0x / TR1
 - std::tr1::array
 - std::tr1::tuple
- Fusion
 - Multiple container types
 - Algorithms!

Applications of Fusion

- Originally developed to do some template "heavy lifting" in Boost.Spirit.
- Applicable to a wide range of generic library implementations – Boost.Spirit, Phoenix, Proto.
- Also suitable for use in application code

Key concepts of Fusion

- Containers
 - vector, list, map....
- Algorithms
 - iteration, querying and transformation
- Iterators
- Views

Fixed at compile time

- The number of types in a container
- The order of types in a container
- Iterator positions

Duality of runtime and compile time

- Fusion sequences are MPL sequences
- MPL sequences are Fusion sequences
- Algorithms, intrinsics etc. have both runtime and compile time forms

```
- fusion::reverse(...) <=>
fusion::result_of::reverse<...>
```

Algorithms are lazily evaluated

- Container independent algorithms
 - push_back, pop_back etc. have one implementation
- Efficient algorithm chaining
 - filter_ifpred>(transform(push_back(cont, v), f()))
 does not lead to excessive copying
- Permits infinite length sequences
 - as long as you don't actually need the whole thing...

A serialization example

- <int>32</int>
- <string>Marsden</string>
- <string>Dan</string>
- <double>123.45</double>

Our example data type

```
struct employee
{
  int age;
  std::string surname;
  std::string forename;
  double salary;
};
```

The conventional approach

Can we eliminate this repetition?

What is it we really intend to do?

"for each member write a formatted XML string to the output stream."

- "for each" sounds like we need to apply some sort of algorithm
- How do we produce a "container" of members to operate upon?

Better expressing our intent (1st Attempt)

```
typedef fusion::vector<int, std::string, std::string, double> employee;

std::ostream& format_employee(std::ostream& stream, employee const& e)
    {
        fusion::for_each(e, format_member(stream));
        return stream;
    }
```

The formatting object

```
struct format_member {
  explicit format_member(std::ostream& stream)
    : stream_(stream) {};
  template<typename T>
  void operator()(T const& val) const
     std::string const type = typeid(val).name();
     stream_ << '<' << type << '>' << val << "</" << type << '>';
  std::ostream& stream_;
```

};

Not for employees only

```
template <typename Sequence>
std::ostream& format_xml(std::ostream& stream, Sequence const& seq)
{
   fusion::for_each(seq, format_member(stream));
   return stream;
}
```

Returning to our original employee

```
std::ostream& format_xml(std::ostream& stream, employee const& e)
{
    return format_xml(
        fusion::vector_tie(e.age, e.surname, e.forename, e.salary));
}
```

The extension alternative

```
BOOST_FUSION_ADAPT_STRUCT(
employee,
(int, age)
(std::string, surname)
(std::string, forename)
(double, salary))
```

Summary so far...

- We can use fusion::vector as a container for a mixture of types
- We can use fusion::vector_tie to produce vectors of references in an ad-hoc manner
- We can use fusion::for_each to apply an operation to every member of a Fusion sequence

A simple bit of formatting

Surname: Marsden

Forename: Dan

Age: 32

Salary: \$ 123.45

The conventional implementation

```
std::ostream& format_employee(std::ostream& s, employee
  const& e)
             << "<b>Surname: " << e.surname << "</b>\n"
  return s
              << "Forename: " << e.forename << "\n"
              << "Age: " << e.years << " years old\n"
              << "Salary: $" << e.salary << '\n';
```

More employees

```
namespace example_lib
                                     struct employee
                                        int age;
  struct employee
                                       std::string surname;
                                       std::string forename;
                                       double salary;
     std::string last_name;
                                     };
     std::string first_name;
     double pay;
     long employee_number;
  };
```

Associating data with keys

```
namespace keys
  struct surname;
  struct forename;
  struct age;
  struct salary;
  struct employee_number;
```

Formatting our original employee

```
std::ostream& format_employee(std::ostream& s,
  example::employee const& e)
  return format(s,
    fusion::map_tie<
      keys::surname, keys::forename,
      keys::age, keys::salary>(
         e.surname, e.forename,
         e.age, e.salary));
```

Formatting the new guy

```
std::ostream& format_employee(std::ostream& s,
  example lib::employee const& e)
  return format(s,
     fusion::map_tie<
        keys::surname, keys::forename,
        keys::employee_number, keys::salary>(
          e.last_name, e.first_name,
          e.employee_number,e.pay));
```

Our flexible formatter version 1.0

```
template<typename Map>
std::ostream& format(std::ostream& s, Map const& m)
  format_optional<keys::surname>(s, m, "<b>Surname: ", "</b>\n");
  format_optional<keys::forename>(s, m, "Forename: ", "\n");
  format_optional<keys::age>(s, m, "Age:", " years\n");
  format_optional<keys::salary>(s, m, "Salary: $", "\n");
  format_optional<keys::employee_number>(s, m, "<b>Employee no:
  ", "</b>\n");
  return s;
```

If the data is there...

```
template < typename Key, typename Map>

typename boost::enable_if < fusion::result_of::has_key < Map, Key>,
    std::ostream&>::type

format_optional(std::ostream& s, const Map& m, const std::string&
    prefix, const std::string& postfix)

{
    return s << prefix << fusion::at_key < Key>(m) << postfix;
}</pre>
```

and if the data is missing...

```
template < typename Key, typename Map>

typename boost::disable_if < fusion::result_of::has_key < Map, Key>,
    std::ostream& >::type

format_optional(std::ostream& s, const Map&, const std::string&, const std::string&)

{
    return s;
}
```

Redundancy again

```
template<typename Map>
std::ostream& format(std::ostream& s, Map const& m)
  format_optional<keys::surname>(s, m, "<b>Surname: ", "</b>\n");
  format_optional<keys::forename>(s, m, "Forename: ", "\n");
  format_optional<keys::age>(s, m, "Age:", " years\n");
  format_optional<keys::salary>(s, m, "Salary: $", "\n");
  format_optional<keys::employee_number>(s, m, "<b>Employee no:
  ", "</b>\n");
  return s;
```

Format specifications

```
template<typename Key>
struct format_spec
  format_spec(std::string const& pre, std::string const& post)
     : pre(pre), post(post)
   {}
   std::string pre;
   std::string post;
};
```

Formatting function object

```
template<typename Map>
struct format_member {
  format_member(Map const& lookup, std::ostream& stream)
     : lookup(lookup), stream(stream) {};
  template<typename Key>
  std::ostream& operator()(format_spec<Key> const& spec) const {
     return format_optional<Key>(stream, lookup, spec.pre, spec.post);
  Map const& lookup;
  std::ostream& stream;
```

Our refactored version

```
template<typename Map>
  std::ostream& format(std::ostream& s, Map const& m) {
     fusion::for_each(
      fusion::make_vector(
       format_spec<keys::surname>("<b>Surname: ", "</b>\n"),
       format_spec<keys::forename>("Forename: ", "\n"),
       format_spec<keys::age>("<b>Age: ", "\n"),
       format_spec<keys::salary>("<b>Salary: $", "\n"),
       format_spec<keys::employee_number>("<b>Employee no: ", "</b>\n")),
      format_member<Map>(m, s));
     return s;
```

The final version

```
template<typename Map, typename FwdSeq>
std::ostream& format(std::ostream& s, Map const& m, FwdSeq
const& formatting)
  fusion::for_each(
    formatting,
    format_member<Map>(m, s));
   return s;
```

Going the extension route again

```
BOOST_FUSION_ADAPT_ASSOC_STRUCT(employee,
    (int, age, keys::age)
    (std::string, surname, keys::surname)
    (std::string, forename, keys::forename)
    (double, salary, keys::salary))
```

Another summary

- A fusion::map allows us to associate values with key types
- We can use fusion::map_tie to form maps in an ad-hoc manner
- Fusion maps provide an alternative to C++ structs to which many types can be adapted

A change to our data format

- <surname>Marsden/surname>
- <forename>Dan</forename>
- <age>32</age>
- <salary>123.45</salary>

The conventional approach

Augmenting the original data

```
std::ostream& format_xml(std::ostream& stream, employee const& e)
  fusion::for_each(
     fusion::zip(
        fusion::vector_tie(e.age, e.surname, e.forename, e.salary),
        fusion::make_vector("age", "surname", "forename", "salary")),
     format_member(stream));
  return stream;
```

Extracting the extra data

```
template<typename T>
  void operator()(T const& t) const
    std::string fieldname = fusion::at_c<1>(t);
    stream << '<' << fieldname << '>'
          << fusion::at_c<0>(t)
          << "</" << fieldname << '>';
```

Yet another summary

- fusion::zip can be used to form a single sequence of tuples from a collection of sequences
- Fusion's transformation algorithms return lazily evaluated views which permits efficient algorithm chaining

Can we improve our syntax

 Formatting arbitrary collections of values is clumsy:

```
format_xml(std::cout, fusion::vector_tie(a, b, c));
format_xml(std::cout, fusion::vector_tie(d, e));
```

Can we make this work like a conventional function?

```
format_xml(std::cout, a, b, c);
format_xml(std:cout, d, e);
```

Formatting function object pt.1

```
struct format_xml_impl
  template<typename Seq>
  struct result
     typedef std::ostream& type;
  };
```

Formatting function object pt.2

```
struct format xml impl {
    template<typename Seq>
    std::ostream& operator()(Seq const& seq) const
       std::ostream& stream = fusion::front(seq);
       boost::fusion::for_each(fusion::pop_front(seq),
 formatter(stream));
       return stream;
```

Variadic formatting

```
fusion::unfused_generic<format_xml_impl> const
format_xml((format_xml_impl()));
...
format_xml(std::cout, 32, 123.45, "marsden", "dan");
format_xml(std::cout, 'a', std::string("cat"), 32000);
```

Summary time again

- Fusion sequences can be parameter lists
- Parameter lists can be Fusion sequences
- Fusion permits implementation of variadic functions until C++0x gives us the real thing
 - Users prefer to work with argument lists
 - Implementation of variadics is easier to achieve with Fusion sequences and algorithms

We've only scratched the surface

- Built in support for standard types
 - boost::array, std::pair, boost::variant, boost::tuple
- Loads more containers, views and algorithms
- Extension support for easy adaptation of your own data types
- Lots lots more....

Thanks!