

# An Introduction to the Rvalue Reference in C++0X

Howard E. Hinnant  
Apple Inc.  
BoostCon 2007

# Outline

- Move Semantics
  - Problem statement
  - Solution
- Perfect Forwarding
  - Problem statement
  - Solution



# Code That We Wish Worked

- We would like to return “big” types by value without worrying if RVO will work or not.
- Functional style programming for expensive value types.
- Workaround: Pass a reference or pointer to the result.

```
vector<int>
compute(...data...);

// Expensive, O(N)!
v = compute(...data...);

void
compute(vector<int>& out,
        ...data ...);
```

# Code That We Wish Worked

Generic, sequence modifying code:

```
Iter i = first;
for (++i; i != last; ++i)
{
    Iter j = i;
    value_type tmp(*j); // Expensive!
    for (Iter k = i; k != first && tmp < *--k; --j)
        *j = *k; // Expensive!
    *j = tmp; // Expensive!
}
```

- Consider when value\_type is vector<string>




# Observation

```
Iter i = first;
for (++i; i != last; ++i)
{
    Iter j = i;
    value_type tmp(*j);
    for (Iter k = i; k != first && tmp < *--k; --j)
        *j = *k;
    *j = tmp;
}
```

- In each of the above expensive copies, the source value is never used again after the copy.

# Observation



- In each of the above expensive copies, the source value is never used again after the copy.



# Observation

```
Iter i = first;
for (++i; i != last; ++i)
{
    Iter j = i;
    value_type tmp(*j); // *j not read again
    for (Iter k = i; k != first && tmp < *--k; --j)
        *j = *k; // *k not read again
    *j = tmp; // tmp not read again
}
```

- In each of the above expensive copies, the source value is never used again after the copy.

# Doesn't Reference Counting Solve this Problem?


- Copy constructing and assignment using reference counting is very fast.
- Just increment/decrement the reference count.
- What could be faster?




# Problems With Reference Counting

- Reference counting creates shared ownership semantics which can be more complicated to reason about than sole-ownership.
- Reference counting generally requires atomic operations in a multi-threaded environment (a potential performance penalty on multiprocessor architectures).

# Shared Ownership vs Sole Ownership



Sole Ownership




Shared Ownership


- Shared ownership is only valuable when items are immutable (const) or when reference semantics is desired (changing one copy updates all copies).
- Cheap assignment and function return is no longer a sufficient reason for the more complex shared ownership model.



# Shared Ownership vs Sole Ownership



Sole Ownership



Shared Ownership

- Shared ownership is only valuable when items are immutable (const) or when reference semantics is desired (changing one copy updates all copies).
- Cheap assignment and function return is no longer a sufficient reason for the more complex shared ownership model.

# vector Move Assignment

(actual syntax)

```
template <class T, class A>
class vector {
public:                                     // handles:
    vector& operator=(const vector& v);    // lvalues
    vector& operator=(vector&& v);        // rvalues
};
```

- The “&&” is an rvalue reference.
- It attracts non-const rvalues when overloaded with const v&.



# What Is An Rvalue Reference

```
class A {};  
A a;  
A&& a_ref1 = A();  
A& a_ref2 = a;
```

- An rvalue reference is just a new type of reference spelled “&&” instead of “&”.

# Why A New Reference Type?


- Use of a new reference type, as opposed to a keyword which means “move” or “rvalue” is desirable because the rvalue reference solves more problems than just move semantics.
  - Move semantics
  - Perfect Forwarding
  - Useful rvalue streams
  - etc.



# What Exactly Is an Rvalue Reference?

- An rvalue reference is exactly like our existing reference (now known as an lvalue reference) with a few exceptions:
  - An rvalue will bind to a non-const rvalue reference.
- Rvalue references and lvalue references are distinct types and overloadable.


# Overloading References



- lvalues can bind to any kind of reference (but prefer lvalue references).
- rvalues can bind to rvalue references or a const-qualified lvalue reference.




# Overloading References



- lvalues can bind to any kind of reference (but prefer lvalue references).
- rvalues can bind to rvalue references or a const-qualified lvalue reference.

# Overloading References



- `const lvalue`s can bind to any const-qualified reference.
- `const rvalue`s can bind to any const-qualified reference.



# Overloading References

```
void f(const A&);    // 1  
void f(A&&);         // 2
```

# Overloading References

```
void f(const A&);    // 1  
void f(A&&);         // 2
```

```
A a;
```

- `f(a);` // 1



# Overloading References

```
void f(const A&);    // 1  
void f(A&&);         // 2
```

```
A a;  
const A ca;
```

- f(a); // 1
- f(ca); // 1

# Overloading References

```
void f(const A&);    // 1  
void f(A&&);         // 2
```

```
A a;  
const A ca;  
const A const_source();
```

- f(a); // 1
- f(ca); // 1
- f(const\_source()); // 1



# Overloading References

```
void f(const A&);    // 1  
void f(A&&);        // 2
```

```
A a;  
const A ca;  
const A const_source();  
A source();
```

- f(a); // 1
- f(ca); // 1
- f(const\_source()); // 1
- f(source()); // 2

# Putting It Together

(clone\_ptr is just an example)

```
template <class T> class clone_ptr {  
private: T* ptr_;  
public: clone_ptr(const clone_ptr& p)  
        : ptr_(p.ptr_ ? p.ptr_->clone() : 0) {}  
clone_ptr& operator=(const clone_ptr& p) {  
    if (this != &p) {  
        T* tmp = p.ptr_ ? p.ptr_->clone() : 0;  
        delete ptr_;  
        ptr_ = tmp;  
    }  
    return *this;  
}  
};
```

- Today's copyable clone\_ptr



# Putting It Together

(clone\_ptr is just an example)

```
template <class T> class clone_ptr {  
private: T* ptr_;  
public:  
    clone_ptr(const clone_ptr& p); // lvalues  
    clone_ptr& operator=(const clone_ptr& p);  
    clone_ptr(clone_ptr&& p)      // rvalues  
        : ptr_(p.ptr_) {p.ptr_ = 0;}  
    clone_ptr& operator=(clone_ptr&& p) {  
        delete ptr_;  
        ptr_ = p.ptr_;  
        p.ptr_ = 0;  
        return *this;  
    }  
};
```

- With move semantics
- no throw
- *fast!*

# The Key to Move Semantics

- Temporaries have no name, and can bind to only one reference.
- If you know a reference refers to a temporary, then you know it is safe to modify that temporary without the rest of your program knowing about it.
- Clients who use `std::move` are making a promise that you can pretend that value is a temporary.



# Example Uses of clone\_ptr

```
typedef clone_ptr<int> CP;  
CP p1;  
CP p2(p1);           // clone_ptr(const clone_ptr&)  
CP p3 = std::move(p2); // clone_ptr(clone_ptr&&)  
CP make_clone_ptr();  // function prototype  
CP p4 = make_clone_ptr(); // clone_ptr(clone_ptr&&)
```

- Client code can copy or move clone\_ptr's around.
- If clone\_ptr's are returned (by value) from a function, the move happens implicitly.
- All std::move does is turn an lvalue into an rvalue (a temporary).

# Writing Move Members

```
class A : B {  
    std::string name_;  
public:  
    A(A&& a)  
        : B(std::move(a)), name_(std::move(a.name_)) {}  
    A& operator=(A&& a)  
        {B::operator=(std::move(a));  
         name_ = std::move(a.name_);  
         return *this;}  
};
```

- When base classes and class data members are involved, delegate move logic with explicit use of `std::move` (just as a rule of thumb).



# Double-move Protection

```
void f1(const A&);  
void f1(A&&);    // can change arg  
void f2(const A&);  
void f2(A&&);    // can change arg  
  
void my_func(A&& a) {  
    f1(a);  
    f2(a);  
}
```

- If `f1()` moves from “a”, `f2()` is likely to have a run-time failure.
- Therefore both calls treat “a” as an lvalue.

# Double-move Protection

```
void f1(const A&);  
void f1(A&&);    // can change arg  
void f2(const A&);  
void f2(A&&);    // can change arg  
  
void my_func(A&& a) {  
    f1(a);           // f1(const A&)  
    f2(a);           // f2(const A&)  
}
```

- If `f1()` moves from “a”, `f2()` is likely to have a run-time failure.
- Therefore both calls treat “a” as an lvalue.



# Double-move Protection

```
void f1(const A&);  
void f1(A&&);    // can change arg  
void f2(const A&);  
void f2(A&&);    // can change arg  
  
void my_func(A&& a) {  
    f1(a);           // f1(const A&)  
    f2(std::move(a)); // f2(A&&)  
}
```

- If a move is desired, it must be explicit.
- This prevents accidental moves.

# Move-Only Types

```
template <class T> class unique_ptr {  
private: T* ptr_;  
    unique_ptr(const unique_ptr& p);  
    unique_ptr& operator=(const unique_ptr& p);  
public:  
    unique_ptr(unique_ptr&& p)  
        : ptr_(p.ptr_) {p.ptr_ = 0;}  
    unique_ptr& operator=(unique_ptr&& p) {  
        delete ptr_;  
        ptr_ = p.ptr_;  
        p.ptr_ = 0;  
        return *this;  
    }  
};
```

- A smart pointer which can *only be* moved, not copied, is a simple variation of clone\_ptr.



# Move-Only Types

```
template <class T> class clone_ptr {  
private: T* ptr_;  
public:  
    clone_ptr(const clone_ptr& p);  
    clone_ptr& operator=(const clone_ptr& p);  
    clone_ptr(clone_ptr&& p)  
        : ptr_(p.ptr_) {p.ptr_ = 0;}  
    clone_ptr& operator=(clone_ptr&& p) {  
        delete ptr_;  
        ptr_ = p.ptr_;  
        p.ptr_ = 0;  
        return *this;  
    }  
};
```

- A smart pointer which can *only be* moved, not copied, is a simple variation of clone\_ptr.

# Move-Only Types

```
template <class T> class unique_ptr {  
private: T* ptr_;  
    unique_ptr(const unique_ptr& p);  
    unique_ptr& operator=(const unique_ptr& p);  
public:  
    unique_ptr(unique_ptr&& p)  
        : ptr_(p.ptr_) {p.ptr_ = 0;}  
    unique_ptr& operator=(unique_ptr&& p) {  
        delete ptr_;  
        ptr_ = p.ptr_;  
        p.ptr_ = 0;  
        return *this;  
    }  
};
```

- A smart pointer which can *only be* moved, not copied, is a simple variation of clone\_ptr.



# Example Uses of unique\_ptr

```
typedef unique_ptr<int> CP;  
CP p1;  
CP p2(p1);           // compile time error  
CP p3 = std::move(p2); // unique_ptr(unique_ptr&&)  
CP make_unique_ptr(); // function prototype  
CP p4 = make_unique_ptr(); // unique_ptr(unique_ptr&&)
```

- Client code can move unique\_ptr's around.
  - If unique\_ptr's are returned (by value) from a function, the move happens implicitly.
  - Clients can not copy unique\_ptr's (compile time error)

# Returning Types From Functions

```
unique_ptr<int> make_unique_ptr()
{
    unique_ptr<int> p(new int);
    *p = 2;
    return p; // Ok, implicit cast to rvalue
} // unique_ptr(unique_ptr&&) called or elided
```

- Where Return Value Optimization is already legal today, there will be an implicit cast to rvalue.
- Thus move-only types are easily returned from factory functions.



# Returning Types From Functions

```
vector<int> make_vector()
{
    vector<int> v(100000);
    v[0] = 2;
    return v;    // Ok, implicit cast to rvalue
}               // vector(vector&&) called or elided
```

- Types that are expensive to copy, but cheap to move, can be efficiently returned from factory functions!
- The move may still be elided via RVO.

# Returning Types From Functions

```
complex<double> make_complex()
{
    complex<double> c(1., 2.);
    c *= c;
    return c; // Ok, implicit cast to rvalue
} // complex(const complex&) called or elided
```

- Types with copy constructors, but without move constructors continue to work exactly as they do today



# Examples Of Move-Only Types

- `fstreams`, `stringstreams`
  - `unique_ptr` (`auto_ptr` replacement)
  - `vector<a move-only type>`
  - `thread` - proposed
  - `lock<Mutex>`
- 
- Any handle-type class which refers to a unique, uncopyable resource, is a good candidate for move-only.

# Containers Of Move-Only Types

```
vector<unique_ptr<int>> vp;  
vp.push_back(unique_ptr<int>(new int(0)));  
*vp.front() = 1;  
vector<ofstream> vf;  
vf.push_back(ofstream("first file"));  
vf[0] << "some data";
```

- Containers may now contain move-only types.
- Elements must be moved in or out, not copied.
- The containers are not copyable, just move-only.
- `vector<unique_ptr<T>>` has the same overhead as `vector<T*>`.



# Unique Ownership Strategies

## Example

- Consider:

```
class Matrix
{
    double* data_;
    int row_;
    int col_;
public:
    Matrix(int row, int col);
    ~Matrix();
    Matrix(const Matrix& m);
    Matrix& operator+=(const Matrix& m);
    ...
};
```

- What is the best way to implement  $\text{Matrix} + \text{Matrix}$ ?

# Simple Unique Ownership Example

- Use Case:

```
Matrix m1 = m2 + m3 + m4 + m5 + m6;
```

- The C++03 solution:

```
Matrix  
operator+(const Matrix& x, const Matrix& y)  
{  
    Matrix r(x);  
    r += y;  
    return r;  
}
```



# Simple Unique Ownership

## Example - Cost

- Cost analysis

Matrix  $m1 = m2 + m3 + m4 + m5 + m6;$

```
allocate row*col doubles for m2+m3 -> t1
allocate row*col doubles for t1+m4 -> t2
allocate row*col doubles for t2+m5 -> t3
allocate row*col doubles for t3+m6 -> m1
deallocate doubles for t3
deallocate doubles for t2
deallocate doubles for t1
```

- 4 allocations, 3 deallocations (assumes RVO).

# Copy On Write - Refcounting

```
class Matrix
{
    MatrixImp* imp_; // refcounted ptr
public:
    Matrix(int row, int col);
    ~Matrix();
    Matrix(const Matrix& m);
    Matrix& operator+=(const Matrix& m);
    ...
};
```

- Assume a quality implementation:
  - Embedded refcount; construction allocates once.
  - Refcount protected by atomics.



# Copy On Write - Refcounting

- Cost analysis

Matrix  $m1 = m2 + m3 + m4 + m5 + m6;$

allocate row\*col doubles for  $m2+m3 \rightarrow t1$

allocate row\*col doubles for  $t1+m4 \rightarrow t2$

allocate row\*col doubles for  $t2+m5 \rightarrow t3$

allocate row\*col doubles for  $t3+m6 \rightarrow m1$

deallocate doubles for  $t3$

deallocate doubles for  $t2$

deallocate doubles for  $t1$

- 4 allocations, 3 deallocations.
- At least 3 atomic operations.

# Why COW Doesn't Help Here

- COW fails to optimize when the reason you're copying is because you want to modify the copy.

```
Matrix
```

```
operator+(const Matrix& x, const Matrix& y)
```

```
{
```

```
    Matrix r(x);    // Lazy copy
```

```
    r += y;         // Copy forced here
```

```
    return r;
```

```
}
```



# Copy On Write - Garbage Collected

```
class Matrix
{
    MatrixImp* imp_; // Let it leak!
public:
    Matrix(int row, int col);
    Matrix(const Matrix& m);
    Matrix& operator+=(const Matrix& m);
    ...
};
```

- Garbage collecting libraries exist for C++.

# Copy On Write - Garbage Collected

- Cost analysis

Matrix  $m1 = m2 + m3 + m4 + m5 + m6;$

allocate row\*col doubles for  $m2+m3 \rightarrow t1$

allocate row\*col doubles for  $t1+m4 \rightarrow t2$

allocate row\*col doubles for  $t2+m5 \rightarrow t3$

allocate row\*col doubles for  $t3+m6 \rightarrow m1$

- 4 allocations.
- 3 allocations likely need be collected later.



# Copy On Write - Garbage Collected

- Garbage collection doesn't correct COW's weakness in the copy/modify pattern.
- Deallocations are not “saved”, they are merely delayed until some time later.

# Alternative GC-Based Solution

- Instead of COW in the Matrix copy constructor, return a leaked Matrix by reference in operator+:

```
Matrix&
operator+(const Matrix& x, const Matrix& y)
{
    Matrix* r = new Matrix(x);
    *r += y;
    return *r;
}
```



# Alternative GC-Based Solution

- Cost analysis

Matrix  $m1 = m2 + m3 + m4 + m5 + m6$ ;

allocate 3 words for  $t1$  in  $m2+m3$

allocate row\*col doubles for  $m2+m3 \rightarrow t1$

allocate 3 words for  $t2$  in  $t1+m4$

allocate row\*col doubles for  $t1+m4 \rightarrow t2$

allocate 3 words for  $t3$  in  $t2+m5$

allocate row\*col doubles for  $t2+m5 \rightarrow t3$

allocate 3 words for  $t4$  in  $t3+m6$

allocate row\*col doubles for  $t3+m6 \rightarrow t4$

allocate row\*col doubles for  $m1$

- 9 allocations.

# Alternative GC-Based Solution

- Sometimes when people know they have garbage collection, they tend to leak with abandon. Often this can degrade performance instead of improve it.
- Garbage collection doesn't mean "never having to worry about memory."



# Rvalue Reference Based Solution

```
Matrix::Matrix(Matrix&& x)
    : data_(x.data_), row_(x.row_), col_(x.col_)
    {x.data_ = 0; x.row_ = x.col_ = 0;}
```

```
Matrix    // this as in C++03
operator+(const Matrix& x, const Matrix& y);
```

```
Matrix
operator+(Matrix&& temp, const Matrix& y) {
    temp += y;
    return std::move(temp);
}
```

- Add a move constructor.
- If op+ sees a temporary, just add to it instead of creating a new temporary.

# Rvalue Reference Based Solution

- Cost analysis

Matrix  $m1 = m2 + m3 + m4 + m5 + m6;$

allocate row\*col doubles for  $m2+m3 \rightarrow m1$



# Rvalue Reference Based Solution

- Cost analysis

Matrix  $m1 = m2 + m3 + m4 + m5 + m6;$

allocate row\*col doubles for  $m2+m3 \rightarrow t1$

$t1$  RVO'd (or moved) out

$t1 += m4$  //  $t1$  moved out

$t1 += m5$  //  $t1$  moved out

$t1 += m6$  //  $t1$  moved into  $m1$

- One allocation.
- Cost analysis does not change if RVO is not applied.

# Cost Summary for Matrix + Matrix

	Allocations	Deallocations	Outstanding
--	-------------	---------------	-------------

- |              |   |   |   |
|--------------|---|---|---|
| • C++03:     | 4 | 3 | 1 |
| • COW (rc)   | 4 | 3 | 1 |
| • COW (gc)   | 4 | 0 | 4 |
| • GC-op+     | 9 | 0 | 9 |
| • Rvalue-ref | 1 | 0 | 1 |
- Move semantics makes unique ownership both simple and practical.



# Cost Summary for Matrix + Matrix


Allocations    Deallocations    Outstanding

- C++03:
- COW (rc)
- COW (gc)
- GC-op+
- Rvalue-ref

4	3	1
4	3	1
4	0	4
9	0	9
1	0	1

- Move semantics makes unique ownership both simple and practical.

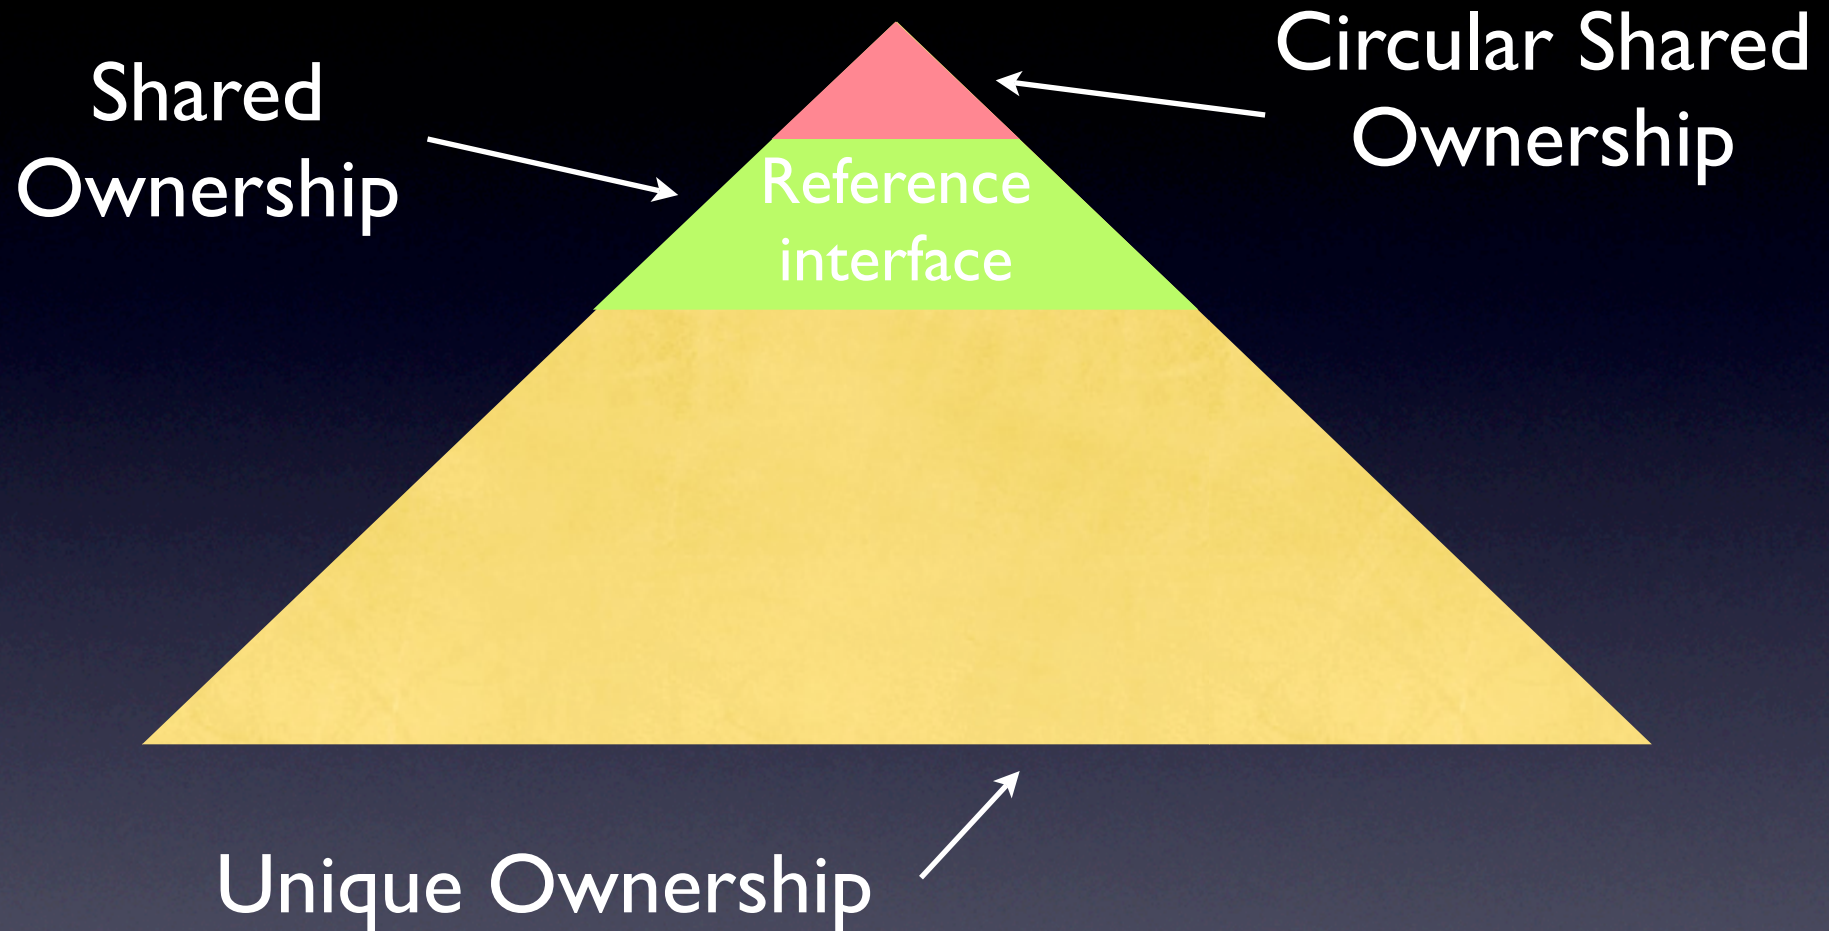
# Memory Ownership Strategies



- Rvalue reference greatly reduces the motivation for shared ownership.



# Memory Ownership Strategies



- Rvalue reference greatly reduces the motivation for shared ownership.

# Generic Algorithm Solution

```
Iter i = first;
for (++i; i != last; ++i)
{
    Iter j = i;
    value_type tmp(std::move(*j));
    for (Iter k = i; k != first && tmp < *--k; --j)
        *j = std::move(*k);
    *j = std::move(tmp);
}
```

- This fact that a value is no longer needed after copied from is communicated with `std::move`.
- For some `value_types` (such as `vector`) this makes the “assignment” very fast.
- For other `value_types` (such as `int`) the `std::move` has absolutely no affect whatsoever.



# Generic Algorithm Solution

```
Iter i = first;
for (++i; i != last; ++i)
{
    Iter j = i;
    value_type tmp(std::move(*j));
    for (Iter k = i; k != first && tmp < *--k; --j)
        *j = std::move(*k);
    *j = std::move(tmp);
}
```

- This fact that a value is no longer needed after copied from is communicated with `std::move`.
- For some `value_types` (such as `vector`) this makes the “assignment” very fast.
- For other `value_types` (such as `int`) the `std::move` has absolutely no affect whatsoever.

# Generic Containers

- I will be using vector (again) just as an example:
  - Move constructor and move assignment make container itself cheaply movable.
  - But what if the container's `value_type` is cheaply moveable?



# Container Moving Interface

```
void push_back(T&& x);  
iterator insert(iterator p, T&& x);
```

- In addition to a move constructor and move assignment operator, a container can *move single* objects into itself.

```
vector<string> v;  
v.push_back("first string"); //moved in  
string s("second string");  
v.push_back(s);              // copied in  
v.push_back(std::move(s));    // moved in
```

# Moving push\_back Implementation

```
void push_back(T&& x) {  
    if (size() < capacity()) {  
        alloc().construct(end(), std::move(x));  
        ...  
    } else  
        ...  
}
```

- Use `std::move` to invoke `x`'s move constructor instead of `x`'s copy constructor.
- If the type does not have a move constructor, its copy constructor is used.



# Moving push\_back Implementation

- There is nothing magic about `allocator::construct`: It simply forwards to placement new, as if:

```
::new (p) T(std::move(x));
```


- If `T` has a move constructor, it is used. Else if `T` has a copy constructor, it is used. Else the placement new will fail to compile.
- All `std::move(x)` does is turn `x` into an rvalue.

# How Much Faster Can Move Be?

- A benchmark has been created which simply constructs, sorts, rotates and destructs a `vector<set<int>>`.
- [http://home.twcny.rr.com/hinnant/cpp\\_extensions/STL\\_benchmarks.html](http://home.twcny.rr.com/hinnant/cpp_extensions/STL_benchmarks.html)
- The benchmark doesn't explicitly contain `std::move` or rvalue reference.
- Run it today and compare your best machine against my 1GHz laptop.



# How Much Faster Can Move Be?



- The difference results from simply avoiding unnecessary trips to the heap.

# Perfect Forwarding

- Consider writing a generic factory function for `shared_ptr<T>`, where `T` is unknown, as are the arguments for constructing `T`.
- Assume for the moment that `T`'s constructor takes a single argument.

```
template <class T, class A1>
shared_ptr<T>
make_ptr(const A1& a1)
{
    return shared_ptr<T>(new T(a1));
}
...
shared_ptr<B> p = make_ptr<B>(2);
```



# Problem using const A1&

```
template <class T, class A1>
shared_ptr<T>
make_ptr(const A1& a1)
{
    return shared_ptr<T>(new T(a1));
}
...
double d = 2;
shared_ptr<B> p = make_ptr<B>(d);
```

- What if B constructs with a non-const reference? B::B(double&)
  - Doesn't compile because make\_ptr uses const A1&.

# Solution: A1 &

```
template <class T, class A1>
shared_ptr<T>
make_ptr(A1& a1)
{
    return shared_ptr<T>(new T(a1));
}
...
double d = 2;
shared_ptr<B> p = make_ptr<B>(d);
```

- Now make\_ptr works with B::B(double&), but...



# Problem using A1&

```
template <class T, class A1>
shared_ptr<T>
make_ptr(A1& a1)
{
    return shared_ptr<T>(new T(a1));
}
...
shared_ptr<B> p = make_ptr<B>(2);
```

- What if B constructs with a double or const double&? B::B(double)
- Doesn't compile because make\_ptr uses A1&. It won't bind to rvalue arguments.

# Lack of Ability to Write Generic Factory Functions

```
make_ptr(A1& a1)  
make_ptr(const A1& a1)
```

- Neither design is sufficiently generic to cover all reasonable use cases.
- We need a way to tell `make_ptr` to accept any lvalue or rvalue, and forward that to `T`'s constructor, with the same cv-qualifications, and lvalue or rvalue status as the original argument supplied to `make_ptr`.



# Rvalue Reference Supplies Perfect Forwarding

```
template <class T, class A1>
shared_ptr<T>
make_ptr(A1&& a1)
{
    return shared_ptr<T>(
        new T(std::forward<A1>(a1)));
}
...
shared_ptr<B> p = make_ptr<B>(2);
```

- A1&& binds to both lvalues and rvalues.
- The cv-qualifications are deduced into the type of A1 (same as today with A1&).

# Rvalue Reference Supplies Perfect Forwarding

```
template <class T, class A1>
shared_ptr<T>
make_ptr(A1&& a1)
{
    return shared_ptr<T>(
        new T(std::forward<A1>(a1)));
}
...
shared_ptr<B> p = make_ptr<B>(2);
```

- Also deduced into the type of A1 is whether an lvalue or rvalue was bound to the parameter.



# Rvalue Reference Supplies Perfect Forwarding

```
template <class T, class A1>
shared_ptr<T>
make_ptr(A1&& a1)
{
    return shared_ptr<T>(
        new T(std::forward<A1>(a1)));
}
...
shared_ptr<B> p = make_ptr<B>(2);
```

- If an lvalue of type B was bound, A1 is B&.
- If an rvalue of type B was bound, A1 is B.

# Rvalue Reference Supplies Perfect Forwarding

```
template <class T, class A1>
shared_ptr<T>
make_ptr(A1&& a1)
{
    return shared_ptr<T>(
        new T(std::forward<A1>(a1)));
}
...
shared_ptr<B> p = make_ptr<B>(2);
```

- `std::forward<B&>` returns an lvalue B.
- `std::forward<B>` returns a rvalue B.



# Rvalue Reference Supplies Perfect Forwarding

```
template <class T, class A1>
shared_ptr<T>
make_ptr(A1&& a1)
{
    return shared_ptr<T>(
        new T(std::forward<A1>(a1)));
}
...
shared_ptr<B> p = make_ptr<B>(2);
```

- Thus the B constructor, called from inside of make\_ptr, sees **exactly** the same argument as supplied from outside of make\_ptr.

# Making it Work With an Arbitrary Number of Arguments

```
template <class T, class ...A>
shared_ptr<T>
make_ptr(A&& a...)
{
    return shared_ptr<T>(
        new T(std::forward<A>(a)...));
}
...
shared_ptr<B1> p1 = make_ptr<B1>(2);
shared_ptr<B2> p2
    = make_ptr<B2>(3, 4.5, std::string("Hi!"));
```

- One simple function perfectly forwards an arbitrary number of arguments!



# Move + Forward

```
struct B2 {  
    B2(int, double, const string&); // copy string in  
    B2(int, double, string&&);      // move string in  
};
```

- Perfect forwarding enables move semantics to work correctly across generic forwarding functions.

# Move + Forward

```
struct B2 {  
    B2(int, double, const string&); // copy string in  
    B2(int, double, string&&);      // move string in  
};
```

- Perfect forwarding enables move semantics to work correctly across generic forwarding functions.

```
// constructs using B2(int, double, string&&)  
shared_ptr<B2> p2  
    = make_ptr<B2>(3, 4.5, string("Hi!"));
```



# So What Does `std::move` Actually Do?

```
template <class T>
inline
typename remove_reference<T>::type&&
move(T&& t)
{
    return t;
}
```

- Accept an rvalue or lvalue and return it as an rvalue.
- The `remove_reference` defeats the unwanted special template deduction for lvalues.

# And std::forward?

```
template <class T>  
struct identity {typedef T type;};
```

```
template <class T>  
inline T&&  
forward(typename identity<T>::type&& t) {  
    return t;  
}
```

- forward<A&> returns A& (an lvalue)
  - References collapse A& && -> A&
- forward<A> returns A&& (an rvalue).



# Summary

- The rvalue reference is a new reference type.
- It solves the problem of move semantics.
- It solves the problem of perfect forwarding.
- Its use is largely hidden at the highest code levels.
  - Use helper functions move and forward.

